

```

# Libraries
%pip install -q seaborn
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
from statsmodels.tsa.stattools import adfuller
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA

# Import Pandas for JupyterLite
import micropip
package_url =
"https://raw.githubusercontent.com/innovationOutside/ouseful\_jupyterlite\_utils/main/ouseful\_jupyterlite\_utils-0.0.1-py3-none-any.whl"
await micropip.install(package_url)
from ouseful_jupyterlite_utils import pandas_utils as pdu

# Import dataset into Pandas dataframe
df = await pdu.read_csv_local("teleco_time_series .csv", "\t")
df[['Day', 'Revenue']] = df['Day,Revenue'].str.split(',', expand =
True)
df.drop('Day,Revenue', axis=1, inplace=True)
df


```

	Day	Revenue
0	1	0
1	2	0.000793191
2	3	0.825541786
3	4	0.32033228
4	5	1.082554085
...
726	727	16.93155866
727	728	17.49066618
728	729	16.80363798
729	730	16.1948135
730	731	16.6207985

```

[731 rows x 2 columns]

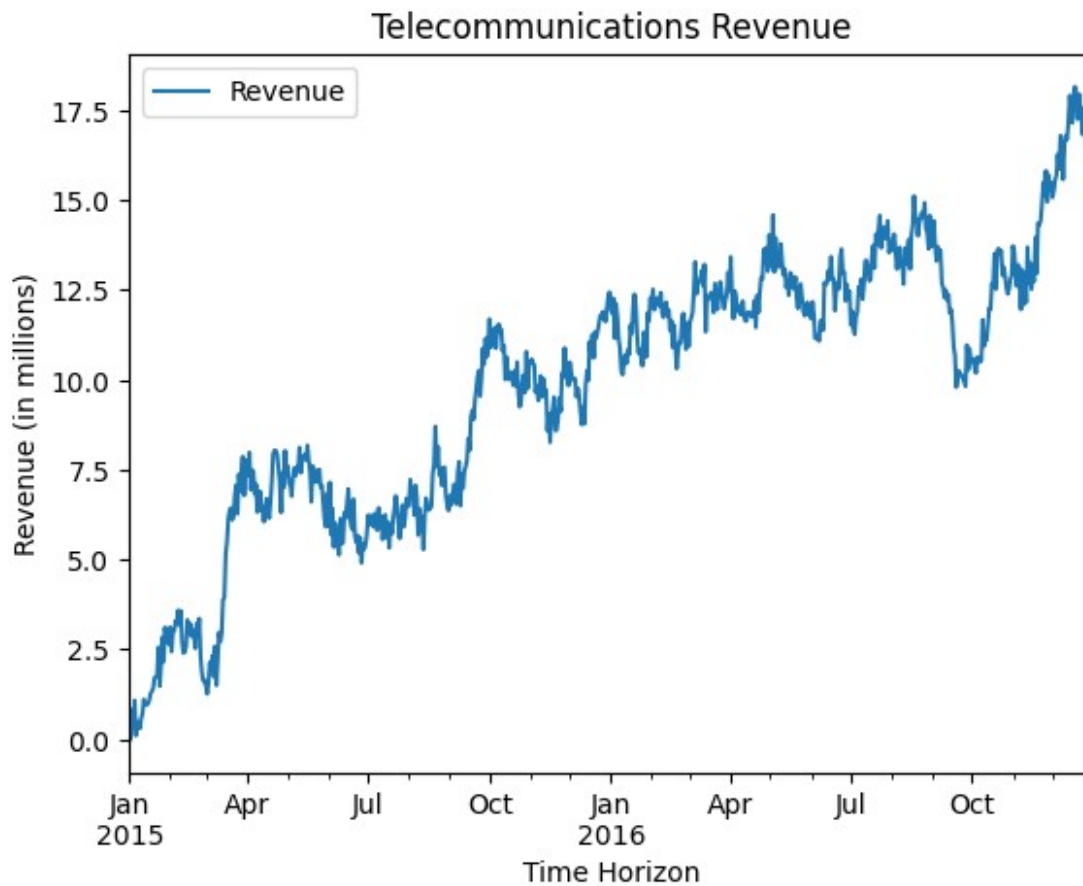
# Convert Revenue column to type float
df['Revenue'] = (df['Revenue']).astype(float)

# Plot line graph
df = df.set_index(pd.date_range(start='2015-1-1', periods=df.shape[0],
freq='D'))

```


```
df.drop('Day', axis=1, inplace=True)
df.plot(title='Telecommunications Revenue', ylabel='Revenue (in
millions)', xlabel='Time Horizon')
```

```
<AxesSubplot:title={'center':'Telecommunications Revenue'},
xlabel='Time Horizon', ylabel='Revenue (in millions)'\>
```



```
# Plot the rolling mean
df.rolling(window = 30).mean().plot()
```

```
<AxesSubplot:>
```




```
# Evaluate the stationarity using Augmented Dickey-Fuller
adf_test = adfuller(df, autolag='AIC')
results = pd.DataFrame(adf_test[:4], columns=['Results '],
dtype=object)
results.index=['Test Statistic', 'p-value', 'Num Lags', 'Num
Observations']
print('Results of Augmented Dickey-Fuller Test for Revenue:\n\n',
results)
```

Results of Augmented Dickey-Fuller Test for Revenue:


	Results
Test Statistic	-1.924612
p-value	0.320573
Num Lags	1.0
Num Observations	729.0

```
# Difference the time series and plot
df_diff = df.diff().dropna()
df_diff.plot(title = 'Differenced (Daily) Telecommunications Revenue',
ylabel='Revenue (in millions)')
```

```
<AxesSubplot:title={'center': 'Differenced (Daily) Telecommunications
Revenue'}, ylabel='Revenue (in millions)'
```



```
# Plot differenced rolling mean  
df_diff.rolling(window = 30).mean().plot()  
<AxesSubplot:>
```



```
# Evaluate stationarity of differenced revenue
adf_diff_test = adfuller(df_diff, autolag='AIC')
results = pd.DataFrame(adf_diff_test[:4], columns=['Results '],
dtype=object)
results.index=['Test Statistic', 'p-value', 'Num Lags', 'Num
Observations']
print('Results of Augmented Dickey-Fuller Test for Differenced
Revenue:\n\n', results)
```

Results of Augmented Dickey-Fuller Test for Differenced Revenue:

	Results
Test Statistic	-44.874527
p-value	0.0
Num Lags	0.0
Num Observations	729.0

```
# Train/Test split
train, test = train_test_split(df_diff, test_size = .2, shuffle =
False, random_state = 123)
train
```

	Revenue
2015-01-02	0.000793
2015-01-03	0.824749

```
2015-01-04 -0.505210
2015-01-05  0.762222
2015-01-06 -0.974900
...
2016-08-03  0.113264
2016-08-04 -0.531705
2016-08-05 -0.437835
2016-08-06  0.422243
2016-08-07  0.179940
```

```
[584 rows x 1 columns]
```

```
test
```

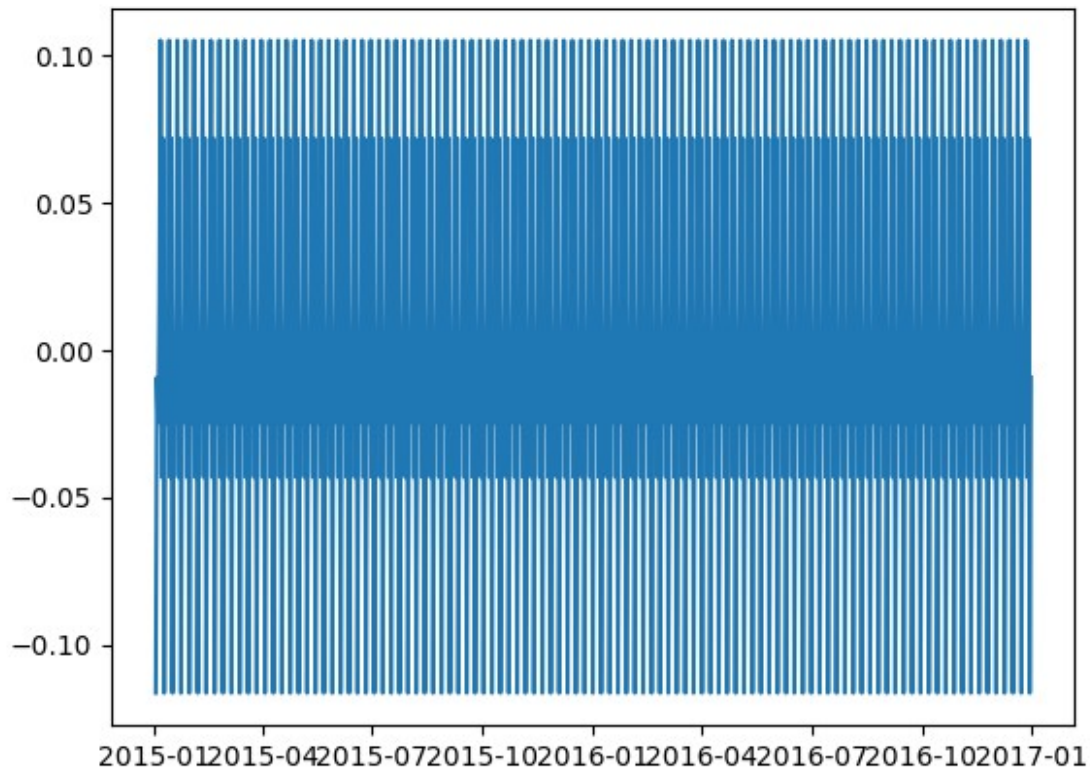
```
Revenue
2016-08-08 -0.531923
2016-08-09  0.157387
2016-08-10 -0.644689
2016-08-11  0.995057
2016-08-12 -0.438775
...
2016-12-27  0.170280
2016-12-28  0.559108
2016-12-29 -0.687028
2016-12-30 -0.608824
2016-12-31  0.425985
```

```
[146 rows x 1 columns]
```

```
# Save training and testing sets to CSV
train.to_csv('D213_task1_train_clean.csv')
test.to_csv('D213_task1_test_clean.csv')
```

```
# Decompose differenced revenue
decomposed_df_diff = seasonal_decompose(df_diff)
plt.plot(decomposed_df_diff.seasonal)
```


```
[<matplotlib.lines.Line2D at 0x7158e80>]
```




```
# Plot seasonal component
plt.plot(decomposed_df_diff.seasonal, marker = 'o')
plt.xlim(pd.to_datetime('2015-01-01'), pd.to_datetime('2015-02-01'))


# Draw red lines at the start of the week
plt.axvline(x=pd.to_datetime('2015-01-05'), color='red')
plt.axvline(x=pd.to_datetime('2015-01-12'), color='red')
plt.axvline(x=pd.to_datetime('2015-01-19'), color='red')
plt.axvline(x=pd.to_datetime('2015-01-26'), color='red')

<matplotlib.lines.Line2D at 0x6eaf110>
```







```
# Plot Autocorrelation and Partial Autocorrelation  
fig, (plot1, plot2) = plt.subplots(1, 2, sharey = True)  
plot_acf(df_diff, lags = 8, zero = False, ax = plot1)  
plot_pacf(df_diff, lags = 8, zero = False, ax = plot2);
```




```
# Plot spectral density  
plt.psd(x = df_diff['Revenue']);
```



```
# Plot decomposed time series  
decomposed_df_diff.plot();
```



```
# Plot residual component  
plt.plot(decomposed_df_diff.resid);
```



```
# ARIMA model
model = ARIMA(train, order = (1,0,0), freq = 'D')
results = model.fit()
print(results.summary())
```

SARIMAX Results

```
=====
=====
Dep. Variable:                Revenue    No. Observations:
584
Model:                        ARIMA(1, 0, 0)    Log Likelihood
-383.946
Date:                        Tue, 29 Aug 2023    AIC
773.893
Time:                        21:32:31    BIC
787.002
Sample:                        01-02-2015    HQIC
779.002
                                - 08-07-2016
```

```
Covariance Type:                opg
```

```
=====
=====
coef    std err          z      P>|z|    [0.025
```


```

0.975]
-----
-----
const          0.0234      0.013      1.758      0.079      -0.003
0.049
ar.L1          -0.4597      0.036     -12.654      0.000      -0.531
-0.388
sigma2         0.2180      0.014      16.034      0.000      0.191
0.245
=====
=====
Ljung-Box (L1) (Q):                0.00   Jarque-Bera (JB):
1.84
Prob(Q):                0.96   Prob(JB):
0.40
Heteroskedasticity (H):            0.97   Skew:
-0.08
Prob(H) (two-sided):            0.83   Kurtosis:
2.77
=====
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients
(complex-step).

# Forecast
forecast = results.get_prediction(start = 584, end = 729, dynamic =
True)
plt.plot(test)
plt.plot(forecast.predicted_mean);

```



```
print(forecast.predicted_mean)
```

```
2016-08-08    -0.048621
2016-08-09     0.056441
2016-08-10     0.008147
2016-08-11     0.030347
2016-08-12     0.020142
```

```
...
2016-12-27     0.023356
2016-12-28     0.023356
2016-12-29     0.023356
2016-12-30     0.023356
2016-12-31     0.023356
```

```
Freq: D, Name: predicted_mean, Length: 146, dtype: float64
```

```
# Make dataframe out of forecast results
```

```
forecast_df = pd.DataFrame(forecast.predicted_mean)
```

```
forecast_df.rename(columns = {'predicted_mean' : 'Revenue'}, inplace =
True)
```

```
forecast_df
```

```
      Revenue
2016-08-08 -0.048621
2016-08-09  0.056441
2016-08-10  0.008147
2016-08-11  0.030347
```

```

2016-08-12    0.020142
...
2016-12-27    0.023356
2016-12-28    0.023356
2016-12-29    0.023356
2016-12-30    0.023356
2016-12-31    0.023356

```

```
[146 rows x 1 columns]
```

```

# Concat a copy of training dataset and copy of forecasted values
train_forecast_df = pd.concat([train.copy(), forecast_df.copy()])

```

```

# Invert differences of daily revenue
train_forecast_df = train_forecast_df.cumsum()
train_forecast_df

```

```

                Revenue
2015-01-02    0.000793
2015-01-03    0.825542
2015-01-04    0.320332
2015-01-05    1.082554
2015-01-06    0.107654
...
2016-12-27    16.952019
2016-12-28    16.975375
2016-12-29    16.998730
2016-12-30    17.022086
2016-12-31    17.045442

```

```
[730 rows x 1 columns]
```

```

# Calculate confidence intervals
conf_int = forecast.conf_int()
conf_int

```

```

                lower Revenue    upper Revenue
2016-08-08    -0.963665         0.866422
2016-08-09    -0.950645         1.063528
2016-08-10    -1.017331         1.033625
2016-08-11    -0.998976         1.059669
2016-08-12    -1.009990         1.050275
...
2016-12-27    -1.006994         1.053705
2016-12-28    -1.006994         1.053705
2016-12-29    -1.006994         1.053705
2016-12-30    -1.006994         1.053705
2016-12-31    -1.006994         1.053705

```

```
[146 rows x 2 columns]
```



```
# Establish dataframe baseline for confidence intervals based on last row of untransformed training set
```

```
baseline = pd.DataFrame({'lower Revenue' : [df['Revenue'][583]],  
    'upper Revenue' : [df['Revenue'][583]], 'date' : ['2016-08-07']})  
baseline['date'] = pd.to_datetime(baseline['date'])  
baseline.set_index('date', inplace = True)  
baseline
```

	lower Revenue	upper Revenue
date		
2016-08-07	13.504886	13.504886

```
# Concat the baseline and confidence intervals
```


```
conf_int = pd.concat([baseline, conf_int])  
conf_int = conf_int.cumsum()  
conf_int = conf_int.loc['2016-08-08' : '2016-12-31']  
conf_int
```

	lower Revenue	upper Revenue
2016-08-08	12.541221	14.371307
2016-08-09	11.590576	15.434835
2016-08-10	10.573245	16.468460
2016-08-11	9.574270	17.528129
2016-08-12	8.564279	18.578404
...
2016-12-27	-129.392812	162.936970
2016-12-28	-130.399806	163.990675
2016-12-29	-131.406800	165.044381
2016-12-30	-132.413794	166.098086
2016-12-31	-133.420788	167.151791

```
[146 rows x 2 columns]
```

```
# Graph predicted and observed data
```

```
plt.title("WGU Telecommunications Daily Revenue, 2015 - 2016")  
plt.xlabel("Date")  
plt.ylabel("Daily Revenue (in Millions USD)")  
plt.plot(train_forecast_df, color = 'green', linestyle = 'dashed')  
plt.plot(df, color = 'blue')  
plt.fill_between(conf_int.index, conf_int['lower Revenue'],  
    conf_int['upper Revenue'], color = 'pink')  
plt.ylim(-2, 20)  
plt.legend(['Predicted', 'Observed'])  
plt.show()
```



```
# Calculate root mean squared error of forecasted vs observed
(untransformed)
rmse = mean_squared_error(df.loc['2016-08-08' : '2016-12-31'],
train_forecast_df.Revenue.loc['2016-08-08' : '2016-12-31'],
squared=False)
print(f"The root mean squared error of this forecasting model is
{round(rmse, 5)}")
```

The root mean squared error of this forecasting model is 2.47394

```
# Plot diagnostics
results.plot_diagnostics();
```

Standardized residual for "R" Histogram plus estimated density

