

D213 Part 2 Performance Assessment Lora Milam

September 25, 2023

```
[2]: # Import Libraries
import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow import keras
import string
from sklearn.feature_extraction.text import CountVectorizer
import random
import nltk
nltk.download('punkt')
nltk.download('stopwords')
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
import string
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.preprocessing.text import Tokenizer
from sklearn import model_selection
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
from keras.callbacks import EarlyStopping
```

```
2023-09-22 07:12:13.841752: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not
find cuda drivers on your machine, GPU will not be used.
2023-09-22 07:12:13.898666: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not
find cuda drivers on your machine, GPU will not be used.
2023-09-22 07:12:13.899786: I tensorflow/core/platform/cpu_feature_guard.cc:182]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 AVX512F FMA, in other operations,
rebuild TensorFlow with the appropriate compiler flags.
2023-09-22 07:12:16.019566: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not
find TensorRT
[nltk_data] Downloading package punkt to
[nltk_data] /home/7d515a27-500e-49a7-accc-
```

```
[nltk_data]      3e4f4339db24/nltk_data...
[nltk_data]      Package punkt is already up-to-date!
[nltk_data]      Downloading package stopwords to
[nltk_data]      /home/7d515a27-500e-49a7-accc-
[nltk_data]      3e4f4339db24/nltk_data...
[nltk_data]      Package stopwords is already up-to-date!
```

```
[3]: # Definition to import dataset as dataframe
def createDF(path, source):
    labels, texts = [], []
    for i, line in enumerate((open(path).read()).split('\n')):
        content = line.split('\t')
        if len(content) > 1:
            texts.append(content[0])
            labels.append(content[1])
    df = pd.DataFrame()
    df['label'] = labels
    df['text'] = texts
    df['source'] = source
    return df
```

```
[4]: # Import datasets into dataframes
amazon_df = createDF('D213 Part 2/amazon_cells_labelled.txt', 'amazon')
imdb_df = createDF('D213 Part 2/imdb_labelled.txt', 'imdb')
yelp_df = createDF('D213 Part 2/yelp_labelled.txt', 'yelp')
```

```
[5]: # Display datasets
display(amazon_df.head())
display(amazon_df.shape)
display(imdb_df.head())
display(imdb_df.shape)
display(yelp_df.head())
display(yelp_df.shape)
```

	label	text	source
0	0	So there is no way for me to plug it in here i...	amazon
1	1	Good case, Excellent value.	amazon
2	1	Great for the jawbone.	amazon
3	0	Tied to charger for conversations lasting more...	amazon
4	1	The mic is great.	amazon

(1000, 3)

	label	text	source
0	0	A very, very, very slow-moving, aimless movie ...	imdb
1	0	Not sure who was more lost - the flat characte...	imdb
2	0	Attempting artiness with black & white and cle...	imdb
3	0	Very little music or anything to speak of.	imdb
4	1	The best scene in the movie was when Gerardo i...	imdb

(1000, 3)

	label	text	source
0	1	Wow... Loved this place.	yelp
1	0	Crust is not good.	yelp
2	0	Not tasty and the texture was just nasty.	yelp
3	1	Stopped by during the late May bank holiday of...	yelp
4	1	The selection on the menu was great and so wer...	yelp

(1000, 3)

```
[6]: # Concatenate datasets
df = pd.concat([amazon_df, imdb_df, yelp_df], ignore_index = True)
```

```
[7]: df
```

```
[7]:
```

	label	text	source
0	0	So there is no way for me to plug it in here i...	amazon
1	1	Good case, Excellent value.	amazon
2	1	Great for the jawbone.	amazon
3	0	Tied to charger for conversations lasting more...	amazon
4	1	The mic is great.	amazon
...
2995	0	I think food should have flavor and texture an...	yelp
2996	0	Appetite instantly gone.	yelp
2997	0	Overall I was not impressed and would not go b...	yelp
2998	0	The whole experience was underwhelming, and I ...	yelp
2999	0	Then, as if I hadn't wasted enough of my life ...	yelp

[3000 rows x 3 columns]

```
[8]: # Describe dataframe
df['chars'] = df.text.apply(len)
df['words'] = df.text.apply(lambda x: len(x.split()))
df['avg_word_len'] = df['chars']/df['words']
df['punctuation'] = df.text.apply(lambda x: len("".join(_ for _ in x if _ in_
↳string.punctuation)))
df['uppercase'] = df.text.apply(lambda x: len([word for word in x.split() if_
↳word.isupper]))
df['titles'] = df.text.apply(lambda x: len([word for word in x.split() if word.
↳istitle]))
df.head()
```

```
[8]:
```

	label	text	source	chars	\
0	0	So there is no way for me to plug it in here i...	amazon	82	
1	1	Good case, Excellent value.	amazon	27	
2	1	Great for the jawbone.	amazon	22	
3	0	Tied to charger for conversations lasting more...	amazon	79	

```
4      1                                The mic is great.  amazon      17
```

	words	avg_word_len	punctuation	uppercase	titles
0	21	3.904762	1	21	21
1	4	6.750000	2	4	4
2	4	5.500000	1	4	4
3	11	7.181818	3	11	11
4	4	4.250000	1	4	4

```
[9]: # Determine if any replicated responses
df['text'].value_counts()
```

```
[9]: Great phone.
2
Not recommended.
2
Works great.
2
I won't be back.
2
I love this place.
2

..
The Songs Were The Best And The Muppets Were So Hilarious.
1
It Was So Cool.
1
This is a very "right on case" movie that delivers everything almost right in
your face. 1
It had some average acting from the main person, and it was a low budget as you
clearly can see. 1
Then, as if I hadn't wasted enough of my life there, they poured salt in the
wound by drawing out the time it took to bring the check. 1
Name: text, Length: 2983, dtype: int64
```

```
[10]: # Determine unique responses
df['text'].unique()
```

```
[10]: array(['So there is no way for me to plug it in here in the US unless I go by a
converter.',
      'Good case, Excellent value.', 'Great for the jawbone.', ...,
      'Overall I was not impressed and would not go back.',
      'The whole experience was underwhelming, and I think we'll just go to
Ninja Sushi next time.',
      'Then, as if I hadn't wasted enough of my life there, they poured salt in
the wound by drawing out the time it took to bring the check.'],
      dtype=object)
```

```
[11]: # Determine min and max length responses
print(min(df['text'].str.len()))
print(max(df['text'].str.len()))
```

7
479

```
[12]: # Determine count of unique words
cnt_vect = CountVectorizer()
features_vect_df = cnt_vect.fit(df.text)
features_df = features_vect_df.get_feature_names_out()
print('df total number of unique words:', len(features_df))
```

df total number of unique words: 5155

```
[13]: # Sample of unique
random.sample(list(features_df), 10)
```

```
[13]: ['commands',
      'ass',
      'securly',
      'few',
      'cuts',
      'documentary',
      'confuses',
      'phones',
      'tremendous',
      'rolled']
```

```
[14]: # Unique word matrix
features_vect_df_trans = cnt_vect.transform(df.text)
type(features_vect_df_trans)
```

```
[14]: scipy.sparse._csr.csr_matrix
```

```
[15]: # Count of matrix values
features_vect_df_trans.getnnz()
```

```
[15]: 31578
```

```
[16]: print("Density of matrix: ", features_vect_df_trans.getnnz()*100/
      ↪(features_vect_df_trans.shape[0]*features_vect_df_trans.shape[1]))
```

Density of matrix: 0.20419010669253151

```
[17]: # Create training Dataframe
train_df = pd.DataFrame(features_vect_df_trans.todense())
train_df.columns = features_df
```

```
train_df[0:1]
```

```
[17]:    00  10  100  11  12  13  15  15g  15pm  17  ...  yucky  yukon  yum  yummy  \
      0   0   0   0   0   0   0   0   0   0   ...   0     0   0     0
      yun  z500a  zero  zillion  zombie  zombiez
      0    0     0    0         0       0       0

[1 rows x 5155 columns]
```


```
[18]: df.text[0:1]
```

```
[18]: 0    So there is no way for me to plug it in here i...
      Name: text, dtype: object
```

```
[19]: # Create dataframe with count
      features_cnt = np.sum(features_vect_df_trans.toarray(), axis = 0)
      features_cnt_df = pd.DataFrame(dict(features_df = features_df, counts =
      ↪features_cnt))
```

```
[20]: # Plot histogram
      plt.figure(figsize = (8,6))
      plt.xlabel('Frequency of Words')
      plt.ylabel('Density')
      plt.hist(features_cnt_df.counts, bins = 10, range = (0, 1000))
```

```
[20]: (array([5.108e+03, 2.700e+01, 8.000e+00, 3.000e+00, 1.000e+00, 1.000e+00,
      3.000e+00, 2.000e+00, 0.000e+00, 0.000e+00]),
      array([ 0., 100., 200., 300., 400., 500., 600., 700., 800.,
      900., 1000.]),
      <BarContainer object of 10 artists>)
```



```
[21]: # Count of words that do not repeat
len(features_cnt_df[features_cnt_df.counts == 1])
```

[21]: 2918


```
[22]: cnt_vect = CountVectorizer(max_features = 3000)
features_vect = cnt_vect.fit(df.text)
features = features_vect.get_feature_names_out()
train_df_features = cnt_vect.transform(df.text)
features_cnts = np.sum(train_df_features.toarray(), axis = 0)
features_cnts = pd.DataFrame(dict(features = features, counts = features_cnts))
features_cnts.sort_values('counts', ascending = False)[0:15]
```

```
[22]:      features  counts
2377      the    1953
85       and    1138
1036     it     789
1033     is     754
2455     to     670
2404    this     643
1334     of     624
```

2814	was	571
1005	in	400
801	for	336
2375	that	316
1315	not	306
2915	with	274
1281	my	254
2736	very	245

```
[23]: # Plot density of positive word
train_df = pd.DataFrame(train_df_features.todense())
train_df.columns = features
train_df['label'] = df.label
sns.barplot(x = 'label', y = 'awesome', data = train_df, estimator = sum)
```

```
[23]: <Axes: xlabel='label', ylabel='awesome'>
```




```
[24]: # Plot density of negative word
train_df = pd.DataFrame(train_df_features.todense())
train_df.columns = features
train_df['label'] = df.label
```



```
sns.barplot(x = 'label', y = 'hate', data =train_df, estimator = sum)
```

```
[24]: <Axes: xlabel='label', ylabel='hate'>
```



```
[25]: df
```

```
[25]:
```

	label	text	source	chars	\
0	0	So there is no way for me to plug it in here i...	amazon	82	
1	1	Good case, Excellent value.	amazon	27	
2	1	Great for the jawbone.	amazon	22	
3	0	Tied to charger for conversations lasting more...	amazon	79	
4	1	The mic is great.	amazon	17	
...
2995	0	I think food should have flavor and texture an...	yelp	66	
2996	0	Appetite instantly gone.	yelp	24	
2997	0	Overall I was not impressed and would not go b...	yelp	50	
2998	0	The whole experience was underwhelming, and I ...	yelp	91	
2999	0	Then, as if I hadn't wasted enough of my life ...	yelp	134	

	words	avg_word_len	punctuation	uppercase	titles
0	21	3.904762	1	21	21

1	4	6.750000	2	4	4
2	4	5.500000	1	4	4
3	11	7.181818	3	11	11
4	4	4.250000	1	4	4
...
2995	12	5.500000	1	12	12
2996	3	8.000000	1	3	3
2997	10	5.000000	1	10	10
2998	16	5.687500	3	16	16
2999	28	4.785714	4	28	28

[3000 rows x 9 columns]

```
[26]: # Review statistics
display(df.groupby(['source', 'label']).describe().loc[:, (slice(None),
↳ ['mean', 'std'])].reset_index())
display(df.groupby(['source', 'label']).describe().loc[:, (slice(None), ['min',
↳ 'max'])].reset_index())
```

	source	label	chars		words		avg_word_len		
			mean	std	mean	std	mean	std	
0	amazon	0	56.824	34.022464	10.578	6.578028	5.653474	1.262250	
1	amazon	1	53.628	35.234764	9.914	6.785772	5.670933	1.113802	
2	imdb	0	77.088	50.887753	13.582	9.036293	5.811760	0.942026	
3	imdb	1	87.456	60.820773	15.128	10.102859	5.826515	0.992945	
4	yelp	0	60.750	34.224935	11.498	6.611916	5.401276	0.896919	
5	yelp	1	55.882	30.228390	10.290	5.831459	5.606414	0.951573	

	punctuation		uppercase		titles	
	mean	std	mean	std	mean	std
0	2.002	1.509488	10.578	6.578028	10.578	6.578028
1	1.842	1.296301	9.914	6.785772	9.914	6.785772
2	2.494	1.954893	13.582	9.036293	13.582	9.036293
3	2.650	2.282525	15.128	10.102859	15.128	10.102859
4	2.000	1.424098	11.498	6.611916	11.498	6.611916
5	1.930	1.658494	10.290	5.831459	10.290	5.831459

	source	label	chars		words		avg_word_len		punctuation		
			min	max	min	max	min	max	min	max	
0	amazon	0	11.0	149.0	1.0	30.0	3.857143	14.0	0.0	11.0	
1	amazon	1	11.0	148.0	1.0	30.0	3.166667	13.0	0.0	9.0	
2	imdb	0	8.0	321.0	1.0	56.0	4.181818	11.5	1.0	14.0	
3	imdb	1	7.0	479.0	1.0	71.0	3.200000	12.0	0.0	18.0	
4	yelp	0	11.0	149.0	2.0	32.0	3.666667	12.5	0.0	11.0	
5	yelp	1	11.0	148.0	1.0	32.0	3.666667	11.0	0.0	19.0	

	uppercase		titles	
	min	max	min	max

```

0      1.0  30.0    1.0  30.0
1      1.0  30.0    1.0  30.0
2      1.0  56.0    1.0  56.0
3      1.0  71.0    1.0  71.0
4      2.0  32.0    2.0  32.0
5      1.0  32.0    1.0  32.0

```

```
[27]: df
```

```

[27]:      label      text  source  chars \
0      0  So there is no way for me to plug it in here i...  amazon    82
1      1      Good case, Excellent value.  amazon    27
2      1      Great for the jawbone.  amazon    22
3      0  Tied to charger for conversations lasting more...  amazon    79
4      1      The mic is great.  amazon    17
...    ...
2995    0  I think food should have flavor and texture an...  yelp    66
2996    0      Appetite instantly gone.  yelp    24
2997    0  Overall I was not impressed and would not go b...  yelp    50
2998    0  The whole experience was underwhelming, and I ...  yelp    91
2999    0  Then, as if I hadn't wasted enough of my life ...  yelp   134

      words  avg_word_len  punctuation  uppercase  titles
0      21      3.904762          1          21      21
1       4      6.750000          2           4       4
2       4      5.500000          1           4       4
3      11      7.181818          3          11      11
4       4      4.250000          1           4       4
...    ...
2995    12      5.500000          1          12      12
2996     3      8.000000          1           3       3
2997    10      5.000000          1          10      10
2998    16      5.687500          3          16      16
2999    28      4.785714          4          28      28

```

```
[3000 rows x 9 columns]
```

```

[28]: # Split the words into tokens
i = 0
df['clean_text'] = ''
for row in df.text:
    # Add spaces after punctuation
    row = row.replace('.', '. ', row.count('.')).replace(',', ', ', row.
↪count(','))
    # Tokenize words
    tokens = word_tokenize(row)
    # Make all words lowercase

```

```

tokens = [token.lower() for token in tokens]
# Remove punctuation
table = str.maketrans('', '', string.punctuation)
# Remove numbers
words = [token.translate(table) for token in tokens]
# Filter stopwords
words = [word for word in words if word.isalnum()]
stop_words = set(stopwords.words('english'))
words = [word for word in words if not word in stop_words]
df['clean_text'][i] = ' '.join(words)
i += 1
df.clean_text = df.source + ' ' + df.clean_text
df.head()

```

/tmp/ipykernel_3674/3478022845.py:19: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['clean_text'][i] = ' '.join(words)

```

[28]:  label                                text  source  chars  \
0      0  So there is no way for me to plug it in here i...  amazon    82
1      1                                Good case, Excellent value.  amazon    27
2      1                                Great for the jawbone.  amazon    22
3      0  Tied to charger for conversations lasting more...  amazon    79
4      1                                The mic is great.  amazon    17

      words  avg_word_len  punctuation  uppercase  titles  \
0      21      3.904762           1           21      21
1       4      6.750000           2            4       4
2       4      5.500000           1            4       4
3      11      7.181818           3           11      11
4       4      4.250000           1            4       4

                                clean_text
0      amazon way plug us unless go converter
1      amazon good case excellent value
2      amazon great jawbone
3  amazon tied charger conversations lasting 45 m...
4      amazon mic great

```

```

[29]: # Add column of count of clean words
df['clean_words'] = df.clean_text.apply(lambda x: len(x.split()))
df.head()

```

```
[29]:
```

	label	text	source	chars	\
0	0	So there is no way for me to plug it in here i...	amazon	82	
1	1	Good case, Excellent value.	amazon	27	
2	1	Great for the jawbone.	amazon	22	
3	0	Tied to charger for conversations lasting more...	amazon	79	
4	1	The mic is great.	amazon	17	

	words	avg_word_len	punctuation	uppercase	titles	\
0	21	3.904762	1	21	21	
1	4	6.750000	2	4	4	
2	4	5.500000	1	4	4	
3	11	7.181818	3	11	11	
4	4	4.250000	1	4	4	

	clean_text	clean_words
0	amazon way plug us unless go converter	7
1	amazon good case excellent value	5
2	amazon great jawbone	3
3	amazon tied charger conversations lasting 45 m...	9
4	amazon mic great	3

```
[30]: print(min(df['clean_words']))
print(max(df['clean_words']))
```

```
2
42
```

```
[31]: features_vect_df2 = cnt_vect.fit(df.clean_text)
features_df2 = features_vect_df2.get_feature_names_out()
print('df2 total number of unique words ', len(features_df2))
```

```
df2 total number of unique words 3000
```

```
[32]: # Describe clean data
df['chars_clean'] = df.clean_text.apply(len)
df['words_clean'] = df.clean_text.apply(lambda x: len(x.split()))
df['avg_word_len_clean'] = df['chars_clean'] / df['words_clean']
df['punctuation_clean'] = df.clean_text.apply(lambda x: len("".join(_ for _ in x
    ↪x if _ in string.punctuation)))
df['uppercase_clean'] = df.clean_text.apply(lambda x: len([word for word in x.
    ↪split() if word.isupper]))
df['titles_clean'] = df.clean_text.apply(lambda x: len([word for word in x.
    ↪split() if word.istitle]))
df.head()
```

```
[32]:
```

	label	text	source	chars	\
0	0	So there is no way for me to plug it in here i...	amazon	82	

1	1	Good case, Excellent value.	amazon	27
2	1	Great for the jawbone.	amazon	22
3	0	Tied to charger for conversations lasting more...	amazon	79
4	1	The mic is great.	amazon	17

	words	avg_word_len	punctuation	uppercase	titles	\
0	21	3.904762	1	21	21	
1	4	6.750000	2	4	4	
2	4	5.500000	1	4	4	
3	11	7.181818	3	11	11	
4	4	4.250000	1	4	4	

	clean_text	clean_words	\
0	amazon way plug us unless go converter	7	
1	amazon good case excellent value	5	
2	amazon great jawbone	3	
3	amazon tied charger conversations lasting 45 m...	9	
4	amazon mic great	3	

	chars_clean	words_clean	avg_word_len_clean	punctuation_clean	\
0	38	7	5.428571	0	
1	32	5	6.400000	0	
2	20	3	6.666667	0	
3	67	9	7.444444	0	
4	16	3	5.333333	0	

	uppercase_clean	titles_clean
0	7	7
1	5	5
2	3	3
3	9	9
4	3	3

```
[33]: display(df.groupby(['source', 'label']).describe().loc[:, (slice(None),
↪ ['mean', 'std'])].reset_index())
display(df.groupby(['source', 'label']).describe().loc[:, (slice(None), ['min',
↪ 'max'])].reset_index())
```

	source	label	chars		words		avg_word_len		\
			mean	std	mean	std	mean	std	
0	amazon	0	56.824	34.022464	10.578	6.578028	5.653474	1.262250	
1	amazon	1	53.628	35.234764	9.914	6.785772	5.670933	1.113802	
2	imdb	0	77.088	50.887753	13.582	9.036293	5.811760	0.942026	
3	imdb	1	87.456	60.820773	15.128	10.102859	5.826515	0.992945	
4	yelp	0	60.750	34.224935	11.498	6.611916	5.401276	0.896919	
5	yelp	1	55.882	30.228390	10.290	5.831459	5.606414	0.951573	

	punctuation		...	words_clean		avg_word_len_clean	\
	mean	std	...	mean	std	mean	
0	2.002	1.509488	...	6.382	3.214458	6.615592	
1	1.842	1.296301	...	6.288	3.382890	6.579993	
2	2.494	1.954893	...	7.974	4.839853	6.389784	
3	2.650	2.282525	...	9.054	5.699830	6.516323	
4	2.000	1.424098	...	6.888	3.415081	6.052106	
5	1.930	1.658494	...	6.444	3.009481	6.252849	

	punctuation_clean		uppercase_clean		titles_clean		\
	std	mean	std	mean	std	mean	
0	1.096950	0.0	0.0	6.382	3.214458	6.382	
1	0.870683	0.0	0.0	6.288	3.382890	6.288	
2	1.033080	0.0	0.0	7.974	4.839853	7.974	
3	1.045467	0.0	0.0	9.054	5.699830	9.054	
4	0.913245	0.0	0.0	6.888	3.415081	6.888	
5	0.927252	0.0	0.0	6.444	3.009481	6.444	

	std
0	3.214458
1	3.382890
2	4.839853
3	5.699830
4	3.415081
5	3.009481

[6 rows x 28 columns]

	source	label	chars		words		avg_word_len		punctuation		\
			min	max	min	max	min	max	min	max	
0	amazon	0	11.0	149.0	1.0	30.0	3.857143	14.0	0.0	11.0	
1	amazon	1	11.0	148.0	1.0	30.0	3.166667	13.0	0.0	9.0	
2	imdb	0	8.0	321.0	1.0	56.0	4.181818	11.5	1.0	14.0	
3	imdb	1	7.0	479.0	1.0	71.0	3.200000	12.0	0.0	18.0	
4	yelp	0	11.0	149.0	2.0	32.0	3.666667	12.5	0.0	11.0	
5	yelp	1	11.0	148.0	1.0	32.0	3.666667	11.0	0.0	19.0	

	...	words_clean		avg_word_len_clean		punctuation_clean		\
	...	min	max	min	max	min	max	
0	...	2.0	17.0	4.333333	10.666667	0.0	0.0	
1	...	2.0	18.0	4.600000	10.000000	0.0	0.0	
2	...	2.0	32.0	4.000000	10.444444	0.0	0.0	
3	...	2.0	42.0	3.666667	10.166667	0.0	0.0	
4	...	2.0	20.0	3.500000	9.666667	0.0	0.0	
5	...	2.0	18.0	3.750000	9.000000	0.0	0.0	

uppercase_clean	titles_clean
-----------------	--------------

	min	max	min	max
0	2.0	17.0	2.0	17.0
1	2.0	18.0	2.0	18.0
2	2.0	32.0	2.0	32.0
3	2.0	42.0	2.0	42.0
4	2.0	20.0	2.0	20.0
5	2.0	18.0	2.0	18.0

[6 rows x 28 columns]

```
[34]: # Pre-padding
tokenizer = Tokenizer(num_words = 50)
seq = tokenizer.texts_to_sequences(df.text)
pad = pad_sequences(seq, padding = "pre", truncating = "pre", maxlen = 42)
pad
```

```
[34]: array([[0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          ...,
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0]], dtype=int32)
```

```
[35]: tokenizer.fit_on_texts(df.clean_text)
word_index = tokenizer.word_index
```

```
[36]: # Post-padding
seq = tokenizer.texts_to_sequences(df.clean_text)
pad = pad_sequences(seq, padding = 'post', truncating = 'post', maxlen = 45)
pad
```

```
[36]: array([[ 1, 46, 27, ..., 0, 0, 0],
          [ 1,  5, 33, ..., 0, 0, 0],
          [ 1,  6,  0, ..., 0, 0, 0],
          ...,
          [ 3, 18, 27, ..., 0, 0, 0],
          [ 3, 43, 27, ..., 0, 0, 0],
          [ 3,  4, 14, ..., 0, 0, 0]], dtype=int32)
```

```
[37]: # Training/testing Split
x_train, x_test, y_train, y_test = model_selection.train_test_split(pad,
↳LabelEncoder().fit_transform(df.label))
```

```
[38]: # Export clean data
df.to_csv('D213_task2_clean')
```



```
[39]: early_stop_monitor = EarlyStopping(monitor = 'val_loss', min_delta = 0,
    ↪patience = 2, verbose = 0, mode = 'auto', baseline = None,
    ↪restore_best_weights = False)
```

```
[40]: # Model hyperparameters
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(3000, 16, input_length = 45),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(6, activation = 'relu'),
    tf.keras.layers.Dense(1, activation = 'sigmoid')
])
```

```
[41]: # Model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 45, 16)	48000
flatten (Flatten)	(None, 720)	0
dense (Dense)	(None, 6)	4326
dense_1 (Dense)	(None, 1)	7
Total params: 52333 (204.43 KB)		
Trainable params: 52333 (204.43 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
[42]: # Compile model
model.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics =
    ↪['accuracy'], run_eagerly = True)
```

```
[43]: x_val = x_train[:5155]
p_x_train = x_train[5155:]
y_val = y_train[:5155]
p_y_train = y_train[5155:]
```

```
[44]: fitmodel = model.fit(x_train, y_train, epochs = 100, batch_size = 128,
    ↪validation_data = (x_val,y_val), verbose = 1, callbacks =
    ↪[early_stop_monitor])
```

Epoch 1/100

18/18 [=====] - 1s 39ms/step - loss: 0.6934 - accuracy:
0.4978 - val_loss: 0.6928 - val_accuracy: 0.5000
Epoch 2/100
18/18 [=====] - 1s 40ms/step - loss: 0.6925 - accuracy:
0.5249 - val_loss: 0.6922 - val_accuracy: 0.5067
Epoch 3/100
18/18 [=====] - 1s 41ms/step - loss: 0.6918 - accuracy:
0.5004 - val_loss: 0.6912 - val_accuracy: 0.5089
Epoch 4/100
18/18 [=====] - 1s 36ms/step - loss: 0.6907 - accuracy:
0.5502 - val_loss: 0.6895 - val_accuracy: 0.5667
Epoch 5/100
18/18 [=====] - 1s 42ms/step - loss: 0.6885 - accuracy:
0.5351 - val_loss: 0.6869 - val_accuracy: 0.5422
Epoch 6/100
18/18 [=====] - 1s 38ms/step - loss: 0.6854 - accuracy:
0.5982 - val_loss: 0.6830 - val_accuracy: 0.5880
Epoch 7/100
18/18 [=====] - 1s 41ms/step - loss: 0.6809 - accuracy:
0.5956 - val_loss: 0.6780 - val_accuracy: 0.5987
Epoch 8/100
18/18 [=====] - 1s 36ms/step - loss: 0.6757 - accuracy:
0.6307 - val_loss: 0.6717 - val_accuracy: 0.6453
Epoch 9/100
18/18 [=====] - 1s 37ms/step - loss: 0.6576 - accuracy:
0.6573 - val_loss: 0.6378 - val_accuracy: 0.6791
Epoch 10/100
18/18 [=====] - 1s 34ms/step - loss: 0.6233 - accuracy:
0.6804 - val_loss: 0.6104 - val_accuracy: 0.6764
Epoch 11/100
18/18 [=====] - 1s 36ms/step - loss: 0.5976 - accuracy:
0.6702 - val_loss: 0.5906 - val_accuracy: 0.6804
Epoch 12/100
18/18 [=====] - 1s 34ms/step - loss: 0.5845 - accuracy:
0.6751 - val_loss: 0.5767 - val_accuracy: 0.6751
Epoch 13/100
18/18 [=====] - 1s 34ms/step - loss: 0.5747 - accuracy:
0.6764 - val_loss: 0.5704 - val_accuracy: 0.6733
Epoch 14/100
18/18 [=====] - 1s 34ms/step - loss: 0.5706 - accuracy:
0.6751 - val_loss: 0.5655 - val_accuracy: 0.6769
Epoch 15/100
18/18 [=====] - 1s 35ms/step - loss: 0.5684 - accuracy:
0.6756 - val_loss: 0.5649 - val_accuracy: 0.6827
Epoch 16/100
18/18 [=====] - 1s 33ms/step - loss: 0.5650 - accuracy:
0.6787 - val_loss: 0.5630 - val_accuracy: 0.6836
Epoch 17/100

```

18/18 [=====] - 1s 34ms/step - loss: 0.5619 - accuracy:
0.6796 - val_loss: 0.5592 - val_accuracy: 0.6791
Epoch 18/100
18/18 [=====] - 1s 35ms/step - loss: 0.5591 - accuracy:
0.6800 - val_loss: 0.5580 - val_accuracy: 0.6858
Epoch 19/100
18/18 [=====] - 1s 35ms/step - loss: 0.5591 - accuracy:
0.6787 - val_loss: 0.5563 - val_accuracy: 0.6849
Epoch 20/100
18/18 [=====] - 1s 35ms/step - loss: 0.5597 - accuracy:
0.6858 - val_loss: 0.5574 - val_accuracy: 0.6787
Epoch 21/100
18/18 [=====] - 1s 34ms/step - loss: 0.5571 - accuracy:
0.6822 - val_loss: 0.5557 - val_accuracy: 0.6813
Epoch 22/100
18/18 [=====] - 1s 38ms/step - loss: 0.5554 - accuracy:
0.6778 - val_loss: 0.5603 - val_accuracy: 0.6840
Epoch 23/100
18/18 [=====] - 1s 39ms/step - loss: 0.5565 - accuracy:
0.6844 - val_loss: 0.5577 - val_accuracy: 0.6769

```

```

[45]: early_stop_monitor = EarlyStopping(monitor = 'val_loss', min_delta = 0,
    ↳patience = 2, verbose = 0, mode = 'auto', baseline = None,
    ↳restore_best_weights = False)
model.fit(x_train, y_train, validation_split = .3, callbacks =
    ↳[early_stop_monitor])

```

```

50/50 [=====] - 2s 31ms/step - loss: 0.5549 - accuracy:
0.6921 - val_loss: 0.5687 - val_accuracy: 0.6593

```

```

[45]: <keras.src.callbacks.History at 0x7fd6ac0508d0>

```

```

[46]: results = model.evaluate(x_test, y_test)

```

```

24/24 [=====] - 0s 11ms/step - loss: 0.5899 - accuracy:
0.6387

```

```

[47]: reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])

```

```

[48]: def decode_review(text):
    return " ".join([reverse_word_index.get(i, '?') for i in text])

```

```

[49]: x_test[1]

```

```

[49]: array([ 3,  4, 16,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
           0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
           0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0], dtype=int32)

```

```
[99]: # Plot Training/Validation Accuracy
sns.set()
acc = fitmodel.history['accuracy']
val = fitmodel.history['val_accuracy']
epochs = range(1, len(acc)+1)

plt.plot(epochs, acc, '-', label = "Training Accuracy")
plt.plot(epochs, val, ':', label = "Validation Accuracy")
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc = 'lower right')
plt.plot()
```

[99]: []



```
[101]: # Save model
model.save('model.h5')
model = keras.models.load_model('model.h5')
model
```

```
/home/7d515a27-500e-49a7-accc-3e4f4339db24/.local/lib/python3.11/site-  
packages/keras/src/engine/training.py:3000: UserWarning: You are saving your  
model as an HDF5 file via `model.save()`. This file format is considered legacy.  
We recommend using instead the native Keras format, e.g.  
`model.save('my_model.keras')`.  
    saving_api.save_model(  

```

```
[101]: <keras.src.engine.sequential.Sequential at 0x7fd6b7a75250>
```

```
[ ]:
```